

# Loom is looming

Rémi Forax / José Paumard



CALVIN & HOBBS © BIL WATTERSON

Don't believe what we are saying !

# What is Loom ?

OpenJDK project started late 2017 by Ron Pressler

Goal: Lowering the cost of concurrency  
should be integrated soon as preview feature

How many threads can I run ?

DEMO

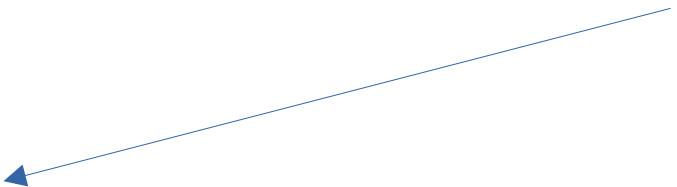
# How many threads can I run ?

```
var barrier = new CyclicBarrier(1_000_001);
for(var i = 0; i < 1_000_000; i++) {
    System.out.println(i);
    new Thread(() -> {
        try {
            barrier.await();
        } catch (InterruptedException | BrokenBarrierException e) {
            throw new AssertionError(e);
        }
    }).start();
}
barrier.await();
```

# How many threads can I run ?

On a MacBook Air M1 (16G of RAM)

```
...  
4063  
4064  
4065  
4066  
[0.373s][warning][os,thread] Failed to start thread "Unknown thread" -  
pthread_create failed (EAGAIN) for attributes: stacksize: 2048k,  
guardsize: 16k, detached.  
[0.373s][warning][os,thread] Failed to start the native thread for  
java.lang.Thread "Thread-4066"  
Exception in thread "main" java.lang.OutOfMemoryError: unable to create  
native thread: possibly out of memory or process/resource limits reached  
    at java.base/java.lang.Thread.start0(Native Method)  
    at java.base/java.lang.Thread.start(Thread.java:1451)  
    at _3_how_many_platform_thread.printHowManyThreads(...java:19)  
    at _3_how_many_platform_thread.main(...java:46)
```



What if I've more than 4066 clients  
for my web server ?



Need to change the model  
1 request  $\Leftrightarrow$  1 threads

# Paradigmatic change

## Asynchronous programming

But I loose the stack trace

=> debugging is harder

=> profiling is harder

=> testing is harder

+

colored function problem

Asynchronous Programming  
is like an addiction

Once you start using it  
It's hard to go back  
... and you loose all your friendly libraries

Solution = coroutine  
(not like Kotlin like golang)

In Java,  
**virtual threads**

DEMO

# Virtual thread

```
// platform threads
var pthread = new Thread(() -> {
    System.out.println("platform " + Thread.currentThread());
});
pthread.start();
pthread.join();

// virtual threads
var vthread = Thread.startVirtualThread(() -> {
    System.out.println("virtual " + Thread.currentThread());
});
vthread.join();
```

# Virtual thread

```
// platform threads  
platform Thread[#14, Thread-0, 5, main]
```

```
// virtual threads  
virtual VirtualThread[#15]/runnable@ForkJoinPool-1-worker-1
```



Use a dedicated thread pool internally



**Warning!** This pool is not the common fork join pool

Or using a *polymorphic* builder



# Thread builder

```
// platform threads
var pthread = Thread.ofPlatform()
    .name("platform-", 0)
    .start(() -> {
        System.out.println("platform " + Thread.currentThread());
    });
pthread.join();

// virtual thread
var vthread = Thread.ofVirtual()
    .name("virtual-", 0)
    .start(() -> {
        System.out.println("virtual " + Thread.currentThread());
    });
vthread.join();
```

How many **virtual** threads can I run ?

DEMO

# How many virtual threads can I run ?

```
var counter = new AtomicInteger();
var threads = IntStream.range(0, 1_000_000)
    .mapToObj(i -> Thread.ofVirtual().unstarted(() -> {
        try {
            Thread.sleep(1_000);
        } catch (InterruptedException e) {
            throw new AssertionError(e);
        }
        counter.incrementAndGet();
    })))
    .toList();

for (var thread : threads) { thread.start(); }
for (var thread : threads) { thread.join(); }
System.out.println(counter); // 1_000_000
```

# Running a thread

Platform (native) thread (starts in **ms**)

- Creates a 2M stack
- System call to ask the OS to schedule the thread

Virtual thread (starts in **μs**)

- Growing stack using stack banging
- Use a specific fork-join pool of pre-created OS threads
  - One OS thread per core

# Concurrency for Loom

Two strategies for concurrency

- Competitive: all threads compete for the CPUs/cores
- Cooperative: each thread hand of the CPUs to the next

Loom does both, carrier threads compete and virtual threads cooperate

How it works under the hood ?

# Uses `jdk.internal.vm.Continuation`

```
var continuation = new Continuation(() -> {  
    System.out.println("C1");  
    Continuation.yield();  
    System.out.println("C2");  
    Continuation.yield();  
    System.out.println("C3");  
});
```

```
System.out.println("start");  
continuation.start();  
System.out.println("came back");  
continuation.start();  
System.out.println("back again");  
continuation.start();  
System.out.println("back again again");
```


Execution:

start  
C1  
came back  
C2  
back again  
C3  
back again again



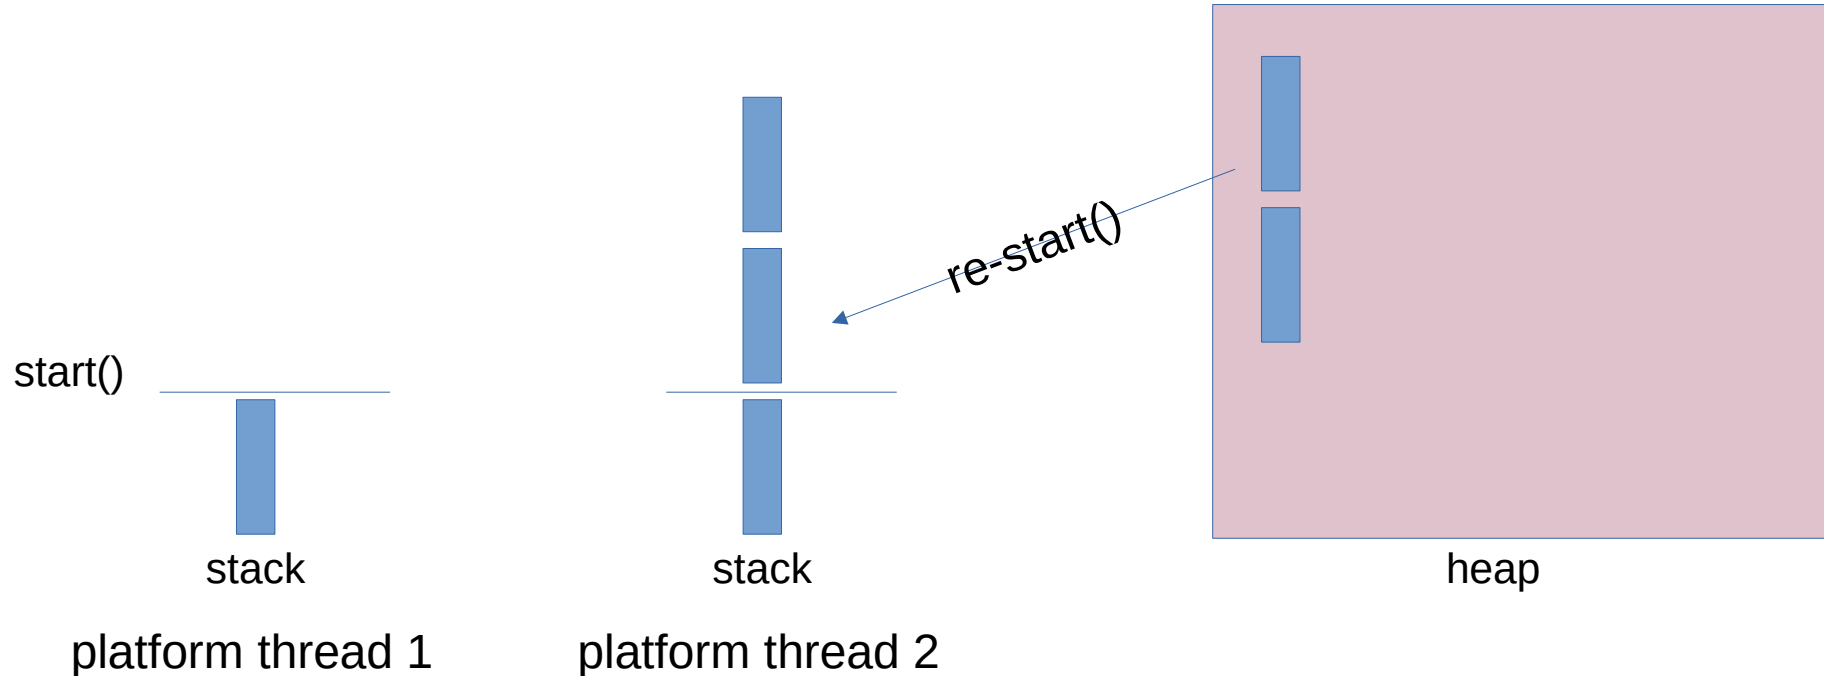
# Continuation.yield()

yield() copy the stack to the heap



# Continuation.start()

start() copy from the heap to another stack



# And in the JDK

All blocking codes are changed to

- Check if current thread is a virtual thread
- If it is, `Continuation.yield()` instead of blocking
- Register a handler that will be called when the OS is ready
  - When the handler is called, find a carrier thread and called `Continuation.start()`

# Example with Thread.sleep()

```
private static void sleepMillis(long millis) throws InterruptedException {  
    Thread thread = currentThread();  
    if (thread instanceof VirtualThread vthread) {  
        long nanos = NANoseconds.convert(millis, MILLISECONDS);  
        vthread.sleepNanos(nanos);  
    } else {  
        sleep0(millis);  
    }  
}
```

```
void sleepNanos(long nanos) throws ...  
    long remainingNanos = ...;  
    while (remainingNanos > 0) {  
        parkNanos(remainingNanos);  
        ...  
    }  
}
```

```
void parkNanos(long nanos) {  
    long startTime = System.nanoTime();  
    boolean yielded;  
    Future<?> unparker = scheduleUnpark(nanos);  
    setState(PARKING);  
    try {  
        yielded = yieldContinuation();  
    } finally {  
        cancel(unparker);  
    }  
  
    // park on the carrier thread for remaining time when pinned  
    if (!yielded) {  
        parkOnCarrierThread(true, deadline - System.nanoTime());  
    }  
}
```

# yield() can fail !

Synchronized block are written in assembly and uses an address on stack

=> the stack frames can not be copied

Native code that does an upcall to Java may use an address on stack

=> the stack frames can not be copied

# *Stealth* rewrite of the JDK for Loom

## Java 13

- JEP 353 Reimplement the Legacy Socket API

## Java 15

- JEP 373 Reimplement the Legacy DatagramSocket API
- JEP 374 Deprecate and Disable Biased Locking

## Java 18

- JEP 416 Reimplement Core Reflection with Method Handles

# Loom: under the Hood

VM creates as many virtual threads as the user want

- It mounts a virtual thread to an available carrier thread when starting
- if blocking, unmount the current virtual thread and mount another virtual thread

# There are still some issues

Synchronized blocks

=> use ReentrantLock instead

Native code that does an upcall

=> no such call in the JDK anymore

Problems with some libraries using native code,  
Hadoop, Spark, ...



# Thread Local issue

1\_000\_000 threads => 1\_000\_000 thread locals ??

# ThreadLocal issue

ThreadLocal implementation store the values in a Map inside `java.lang.Thread`

- Does not scale well !

=> provide a more lightweight implementation

`jdk.incubator.concurrent.ScopeLocal`

DEMO

# Thread Local

```
private static final ThreadLocal<String> USER = new ThreadLocal<>();

private static void sayHello() {
    System.out.println("Hello " + USER.get());
}

public static void main(String[] args) throws InterruptedException {
    var vthread = Thread.ofVirtual()
        .allowSetThreadLocals(true)
        .start(() -> {
            USER.set("Bob");
            try {
                sayHello();
            } finally {
                USER.remove();
            }
        });
    vthread.join();
}
```

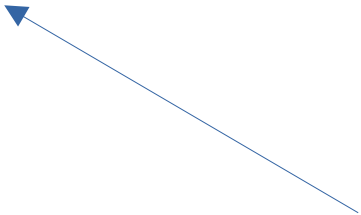
Can be used to disallow thread locals  
throw an ISE when calling ThreadLocal.set()

# Scope Local

```
private static final ScopeLocal<String> USER = ScopeLocal.newInstance();

private static void sayHello() {
    System.out.println("Hello " + USER.get());
}

public static void main(String[] args) throws InterruptedException {
    var vthread = Thread.ofVirtual()
        .allowSetThreadLocals(false)
        .start(() -> {
            ScopeLocal.where(USER, "Bob", () -> {
                sayHello();
            });
        });
    vthread.join();
}
```



Assign the value for the scope

# ScopeLocal

*WARNING API in progress*

- Replacement for ThreadLocal
- Stores the value inside the stack, not inside `java.lang.Thread`
  - => faster (if not too many locals)
  - => use far less memory
- API amenable to JITs

# Executor and structured concurrency

# Executors

An executor recycle the threads

- Do we need an executor if creating a virtual thread does not cost much ?

An executor as another role

- Manage all the submitted task
  - But cancellation/exception management is wrong !



*WARNING API in progress*

# Structured Concurrency

Use syntactic constructions to represent the  
dependency tree of the tasks

DEMO

# VirtualThreadPerTaskExecutor

```
var executor = Executors.newCachedThreadPool();  
//var executor = Executors.newVirtualThreadPerTaskExecutor();
```

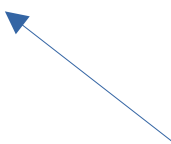
```
var future1 = executor.submit(() -> {  
    Thread.sleep(10);  
    return 42;  
});
```

```
var future2 = executor.submit(() -> {  
    Thread.sleep(1_000);  
    return 100;  
});
```

```
executor.shutdown();
```

```
var result = future1.get() + future2.get();  
System.out.println(result);
```

```
// everything is fine here, right !
```



Special executor using virtual threads  
Warning! the carrier threads are *daemon*

# Running task (1) with an Executor


```
var executor = Executors.newCachedThreadPool();

var future1 = executor.submit(() -> {
    Thread.sleep(10);
    return 42;
});

var future2 = executor.submit(() -> {
    Thread.sleep(1_000);
    System.out.println("end");
    return 100;
});

executor.shutdown();

//var result = future1.get() + future2.get();
var result = future1.get();
System.out.println(result);

// future2 still running here !  Oops !
```

# Running task (2) with an Executor

```
var executor = Executors.newCachedThreadPool();

var future1 = executor.<Integer>submit(() -> {
    throw new AssertionError("oops");
});

var future2 = executor.submit(() -> {
    Thread.sleep(1_000);
    System.out.println("end");
    return 100;
});

executor.shutdown();

Try {
    var result = future1.get() + future2.get();
    System.out.println(result);
} catch (ExecutionException e) {
    throw new AssertionError(e.getCause());
}
```

*// future2 still running here !*

←————— Oops !


# jdk.incubator.concurrent.StructuredTaskScope

```
try (var scope = new StructuredTaskScope<>()) {  
    var start = System.currentTimeMillis();  
  
    var future1 = scope.fork(() -> {  
        Thread.sleep(1_000);  
        return 42;  
    });  
  
    var future2 = scope.fork(() -> {  
        Thread.sleep(1_000);  
        return 100;  
    });  
  
    scope.join();  
  
    var end = System.currentTimeMillis();  
    System.out.println("elapsed " + (end - start));  
    var result = future1.resultNow() + future2.resultNow();  
    System.out.println(result);  
} // call close() !
```

Wait for all computations



Throw an exception in case of dangling tasks



# Future state()

```
try (var scope = new StructuredTaskScope<>()) {  
    var future = scope.fork(() -> {  
        Thread.sleep(1_000);  
        return 42;  
    });  
  
    System.out.println(future.state()); // RUNNING  
  
    //scope.shutdown();  
  
    scope.join();  
  
    System.out.println(future.state()); // SUCCESS  
}
```

# Future state() with a shutdown()

```
try (var scope = new StructuredTaskScope<>()) {  
    var future = scope.fork(() -> {  
        Thread.sleep(1_000);  
        return 42;  
    });  
  
    System.out.println(future.state()); // RUNNING  
  
    scope.shutdown();  
  
    scope.join();  
  
    System.out.println(future.state()); // CANCEL  
}
```



# StructuredTaskScope.shutdown()

If called by the main thread


- Shutdown all the tasks

If called by one task

- Shutdown all the tasks but the caller task

# Shutdown() on failure

```
try (var scope = new StructuredTaskScope.ShutdownOnFailure()) {  
    var start = System.currentTimeMillis();  
  
    var future1 = scope.<Integer>fork(() -> {  
        throw new AssertionError("oops");  
    });  
  
    var future2 = scope.fork(() -> {  
        Thread.sleep(1_000);  
        System.out.println("end");  
        return 42;  
    });  
  
    scope.join();  
  
    var end = System.currentTimeMillis();  
    System.out.println("elapsed " + (end - start));  
    var result = future1.resultNow() + future2.resultNow();  
    System.out.println(result);  
}  
// future and future2 are not running here !
```



This task fails

# Shutdown() on success

```
try (var scope = new StructuredTaskScope.ShutdownOnSuccess<Integer>()) {  
    var start = System.currentTimeMillis();  
  
    var future1 = scope.fork(() -> {  
        Thread.sleep(1_000);  
        return 42;  
    });  
  
    var future2 = scope.fork(() -> {  
        Thread.sleep(42);  
        return 100;  
    });  
  
    scope.join();  
  
    var end = System.currentTimeMillis();  
    System.out.println("elapsed " + (end - start));  
    //System.out.println(future1.resultNow());  
    //System.out.println(future2.resultNow());  
    System.out.println(scope.result());  
}
```

← This task completes first

← The result is stored in the scope

# Summary

# Loom

Will be integrated soon Java 19 (or 20)

Introduce virtual threads (better for latency)

Synchronized() block and native code with an upcall make the code slower

More APIs to come (`jdk.incubator.concurrent`)

- ScopeLocal
- StructuredTaskScope

# Questions ?

<https://github.com/forax/loom-fiber>